

Backup and Change Management

Security Design

System Examination

System Configuration

Firewalls and Filters

Hardening Software

Backups and Change Management

Access Control and Authentication

Virtual Private Networking

Logging and Monitoring

Security Policy and Management Support

SELinux

Linux Systems Security

Why Backups Matter for Security

Data Protection: Safeguard data from loss, corruption, or theft

Disaster Recovery: Essential for system recovery after incidents

Risk Mitigation: Provides an effective response against hardware failure, cyber threats, and human error

Security Assurance: Backups add a layer of confidence in system resilience

Types of Backups

Full Backup: Entire system or specific partitions are copied, ideal for first-time or periodic backups

Incremental Backup: Only changes since the last backup are saved, reducing storage and time

Differential Backup: Backs up changes since the last full backup, faster than full but requires more storage than incremental

Snapshot Backup: Captures the system state at a specific point in time, useful for quick restoration

Essential Tools for Linux Backups

rsync: Efficient file syncing and copying, ideal for both local and remote backups

tar: Archiving tool that compresses files into a single package, useful for full backups

dd: Low-level copying for disk and partition images, often used for system cloning

Timeshift: Snapshot tool, commonly used in desktop environments for system restoration

Bacula and Amanda: Open-source backup software for managing large-scale, enterprise-level backups

Legacy Tools



https://en.wikipedia.org/wiki/Manuscript_culture

- `cp` is the original way to make a copy of files, but assumptions people make cause problems in using it
- GUI-based drag and drop tools make the assumptions problem worse
- `cpio`, `tar` are archival tools created to copy files to backup media (tape by default) - satisfactory for years but they make it cumbersome to manage backup media
- Various software packages provide frontends to these tools in order to make backup/restore easier to manage and more robust

Rsync

- **Features of Rsync:**

- Resembles a smart cp
 - Can preserve special files such as links and devices
 - Can copy to or from local or remote destinations
 - Provides standby device capability by synchronizing data stores
 - Can compress data in transit and only sends changed data
 - Can use ssh to encrypt remote backup transfers without requiring encryption of backup storage
 - Light on bandwidth, heavier on cpu and memory
- **Agent based:** Requires rsync be available at both ends of the connection if used over the network
 - Can be used with network shares to avoid requirement for rsync at both ends of connections at a price in performance and network utilization

Rsync Examples

- **Simple local archive of user data**

```
rsync -ahv --delete ~/* /backup/username  
rsync -ahv --delete --exclude=~ /myBackup ~/* ~ /myBackup
```

- **Local archive with history**

```
date=$(date "+%F-%H-%M")  
rsync -ahv --link-dest=/backup/latest --delete --exclude={"/proc/*,/tmp/*,/run/*,/dev/*,/sys/*,/mnt/*,  
*/lost+found,/media/*,/backup} /* /backup/$date &&  
ln -nsf /backup/$date /backup/latest
```

- **Remote archive with history**

```
date=$(date "+%F-%H-%M")  
rsync -ahv --link-dest=/backup/latest --delete --exclude={"/proc/*,/tmp/*,/run/*,/dev/*,/sys/*,/mnt/*,  
*/lost+found,/media/*,/backup}/* rsync-user@host:/backup/$date &&  
ssh rsync-user@host ln -nsf /backup/$date /backup/latest
```

Rsync Restore

- Restoration of user data can be done using rsync command to copy data from backup location
- Can be copied to user staging or live data location
- Since the backup filesystem is a synchronized duplicate of the original filesystem, normal tools can be used for examining what is available in a backup
- System recovery requires additional tasks to be performed to deal with the boot block
 - Rsync the backup to a new drive
 - Install the boot block using GRUB on the new drive
 - Install the new drive and boot from it

Duplicity

- Like rsync but does not require duplicity on the receiving backup host, stores metadata with backup
- Can use many different types of backend storage, including Amazon S3, although they require additional backend packages and configuration effort
- Encrypts tarballs as a storage mechanism, stores deltas in separate tarballs, may have more steps involved in recovery
- Written in python, heavy on bandwidth, lighter on cpu and memory, reimplements ssh in python
- `duplicity /what/to/backup sftp://user@host/backup/directory`

Other Backup Solutions

- **GUI Frontends:** Many GUI frontends to the rsync, duplicity, and tar cli tools are available
 - Backintime, DejaDup, CronoPete, Timeshift, Duplicati
 - Bacula, Amanda
 - Mondo Rescue for disaster recovery

Backup Strategies and Planning

Determine Critical Data: Identify which data or configurations are essential for your system's function and security

Frequency and Scheduling: Decide on a backup schedule—daily, weekly, or monthly—based on the rate of data change

Storage Location: Choose secure storage options—local external drives, network-attached storage (NAS), or cloud

Redundancy: Use the **3-2-1 Rule**: Three copies of data, on two different storage types, with one offsite

Automating Backups with Linux

Cron Jobs: Schedule regular backups by creating cron jobs to automate backup tools like rsync or tar

Shell Scripts: Customize backup scripts to automate complex backup processes, including notifications

Monitoring Tools: Implement tools like **Nagios** or **Zabbix** to monitor backup completion and system status

Secure Backup Practices

Encryption: Use tools like gpg to encrypt backups, especially when storing offsite or in the cloud

Access Control: Limit access to backup files and directories to minimize the risk of unauthorized access

Integrity Checks: Regularly verify backups with checksums or hashes to ensure data integrity

Versioning: Keep different versions of backups for flexibility in recovery, especially for incremental and differential backups

Testing Backups for Reliability

Testing Frequency: Regularly restore data to verify backup integrity and usability

Testing Methods: Test in a controlled environment to avoid impacting live systems

Logging and Auditing: Maintain logs of backup tests and results to identify and resolve issues quickly

Backup Challenges

Storage Constraints: Balancing storage costs with the need for redundancy and retention policies

Data Sensitivity: Managing encryption and access controls for sensitive data backups

Complex Environments: Handling mixed environments with multiple Linux distributions and configurations

Human Error: Mitigating risks associated with misconfiguration or accidental data deletion

Conclusion

Key Takeaways:

- Regular backups are critical for Linux system resilience
- Automating and securing backups strengthens data security
- Regular testing ensures backups are reliable when needed

Final Thought: Effective backup planning is a foundational aspect of Linux system security

Change Management

What is Change Management?

- **Definition:** A structured approach to managing alterations in IT environments, ensuring changes are efficiently and securely implemented.
- **Objectives:** Maintain system stability, ensure security, and minimize service disruptions.
- **Relevance to Linux Systems:** With Linux's flexibility and potentially frequent updates, change management is critical for both enterprise and individual users to maintain system security.

Why Change Management is Essential for Security

Risk Mitigation: Reduces the likelihood of unauthorized or unexpected changes that can introduce vulnerabilities.

System Integrity: Prevents untested changes from disrupting core system functions.

Compliance: Helps meet regulatory standards and organizational policies by documenting changes and approvals.

Incident Response: Allows for quick identification of recent changes that could be linked to new issues or vulnerabilities.

Types of Changes

Software Updates: Regular updates to the OS, applications, and security patches.

Configuration Changes: Adjustments to settings and parameters that affect performance or security (e.g., firewall rules, network configurations).

Hardware Changes: Adding or replacing hardware components that require new drivers or configurations.

System Upgrades: Major updates or migration to new Linux distributions, which can impact dependencies and functionality.

Key Steps in Change Management

Request: Initiate a formal request for change, including a description, reason, and priority.

Review and Approval: Evaluate the change for potential impact and obtain approvals from relevant stakeholders.

Testing: Test the change in a staging environment to identify any issues before implementation.

Implementation: Execute the change with proper monitoring.

Documentation: Record details of the change, reasons, testing results, and approval history.

Post-Change Review: Assess the change's success and impact, addressing any issues that arise.

Tools for Change Management

Git: Version control for configuration files and scripts, enabling rollback to previous configurations.

Ansible/Puppet/Chef: Configuration management tools that automate the deployment of consistent settings across multiple systems.

Auditd: Monitors and logs changes to critical files and configurations.

Syslog: Centralized logging to capture system events, helping track when and how changes were made.

Redmine or Jira: Project management and ticketing systems to track and document change requests and statuses.

Automating Change Management

Automated Testing: Use CI/CD pipelines to automatically test changes in a controlled environment.

Scheduled Updates: Automate updates and patches during low-impact times to minimize disruptions.

Rollback Mechanisms: Set up snapshots or backup configurations to quickly revert changes if issues occur.

Notifications and Alerts: Configure alerts to notify admins of unauthorized or unexpected changes.

Change Management Best Practices

Change Advisory Board (CAB): In larger environments, establish a board to review and approve changes.

Define Clear Policies: Set policies for which changes require approval, testing, and documentation.

Minimal Downtime: Plan implementations to minimize impact on users, utilizing rolling updates or blue-green deployments if possible.

Regular Audits: Conduct audits to ensure that changes comply with policies and don't introduce vulnerabilities.

Challenges of Change Management in Linux Systems

Configuration Drift: Unplanned changes can lead to inconsistencies, particularly across multiple servers.

Human Error: Risk of misconfigurations that can impact security and system stability.

Resource Intensive: Requires time and resources for proper testing and documentation, which can be challenging in low investment environments.

Interdependencies: Complex dependencies between software packages and configurations can complicate changes.

Conclusion

Key Takeaways:

- Structured change management reduces security risks and improves system stability.
- Automated tools and consistent policies streamline the change process.
- Documentation and testing are crucial to prevent and quickly address issues.

Final Thought: Proactive change management is vital for reliable, secure systems.